

# CSI 410. Database Systems I – Spring 2021

## Programming Assignment II

The total grade for this assignment is 100 points. The deadline for this assignment is **11:59 PM, May 10, 2021**. *Submissions after this deadline will not be accepted.* Students are required to enter the UAlbany Blackboard system and then upload a .zip file (in the form of [first name]\_[last name].zip) that contains the Eclipse project directory and a short document describing:

- any missing or incomplete elements of the code
- any changes made to the original API
- the amount of time spent for this assignment
- suggestions or comments if any

No submission of the above document will lead to a *loss of 5 grade points*. Also, please add *comments* in your code. Code with (almost) no comments may result in a loss of 1 grade point per method to implement.

---

This programming assignment focuses on implementing two relational operators (selection and aggregation). You first need to run Eclipse on your machine and import the “hdb\_query” project (see Appendix A). This assignment requires the “hdb\_data” project that you worked on in the past. You should be able to do this assignment as long as `Tuple(RelationSchema, Object...)` in `Tuple.java` and `attributeIndex(String)` and `attributeType(int)` in `RelationSchema.java` are correctly implemented (other parts in “hdb\_data” are unlikely to affect this project). Please generate a Javadoc API document and then take a look at the document as well as the source code to familiarize yourself with this assignment. For this assignment, we have provided you with a set of classes (in particular, see `ProjectionOperator` which will help you understand how a relational operator can be implemented as well as `TupleArrayReader`, `SelectionOperator`, `AggregateOperator`, and `Aggregator` that you need to complete by adding more code). Your code will be graded by running a set of unit tests and then examining your code. Passing unit tests does not necessarily guarantee that your implementation is correct and efficient. Please strive to write correct and efficient code. If you have questions, please contact the TA(s) or the instructor. The remainder of this document describes the components that you need to implement.

### Part 1. Operators and TupleArrayReader (30 points)

In this assignment, all of the relational operators support the following four methods (see `Operator.java`):

- `hasNext()`: determines whether or not the operator has the next output `Tuple`.
- `next()`: returns the next output `Tuple`.
- `rewind()`: rewinds the operator in order to retrieve all of the output `Tuples` again.
- `outputSchema()`: returns the output schema of the operator.

It should be noted that the `Operator` interface inherits the `hasNext()` and `next()` methods from the standard Java `Iterator` interface.

In this part, you need to complete the code in `TupleArrayReader.java`. Each `TupleArrayReader` outputs, given an array of `Tuples`, the `Tuples` in that array. The methods to complete are `hasNext()`, `next()`, and `rewind()`. Implement these methods using the `currentIndex` member variable which is to remember the index of the next `Tuple` to retrieve/return during each iteration over the `Tuples`. When all of the above methods are implemented correctly, your code will pass the unit tests in `TupleArrayReaderTest.java`.

## Part 2. Selection (50 points)

In this part, you need to complete the code in `SelectionOperator.java`. For this task, it might be helpful to understand the implementation of `ProjectionOperator`. Each `SelectionOperator` outputs, given a series of `Tuples` from another operator, the `Tuples` that satisfy a predicate given to the `SelectionOperator` when it was constructed. In other words, it filters out all `Tuples` that do not match its predicate. Each `SelectionOperator` has its own `ExpressionEvaluator` which can evaluate an expression (i.e., the predicate) for each input `Tuple`. Given `ExpressionEvaluator evaluator` and an input `Tuple t`, (`evaluator.evaluate(t) == Boolean.TRUE`) indicates that `t` satisfies the predicate that `evaluator` uses (i.e., the `SelectionOperator` must output `t`). The `SelectionOperator` can also be viewed as an iterator over all of the `Tuples` that it outputs. Your implementation should not keep all of the output `Tuples` in the memory since it may be infeasible in practical situations (i.e., the memory may not be able to preserve all of the output `Tuples` when a large number of the `Tuples` are output). Instead, the `SelectionOperator` (which is also viewed as an iterator) needs to find, whenever the `hasNext()` and `next()` methods are called, the next input `Tuple` that satisfies the predicate (on-demand, pull-based pipeline; see Section 15.7 in the textbook).

The constructor/methods to complete are as follows:

- `SelectionOperator(Operator input, String predicate)`: constructs a `SelectionOperator`. This constructor needs to create an `ExpressionEvaluator` for evaluating the predicate on each input `Tuple` and may do some additional work (depending on your implementation) to support the `hasNext()` and `next()` methods.
- `outputSchema()`: returns the output schema of the `SelectionOperator`. This output schema is the same as the input schema of the `SelectionOperator`. Consider using a method or a member variable provided by a super-type of `SelectionOperator` to get the input schema of the `SelectionOperator`.
- `hasNext()`: determines whether or not the `SelectionOperator` has the next output `Tuple`.
- `next()`: returns the next output `Tuple`.
- `rewind()`: rewinds the operator in order to retrieve all of the output `Tuples` again.

When all of the above methods are implemented correctly, your code will pass the unit tests in `SelectionOperatorTest`.

## Part 3. Aggregation: `AggregateOperator.java` (5 points)

In this part, you need to complete `AggregateOperator.java`. An `Aggregator` groups all `Tuples` (by user-specified attributes) from an input `Operator` and outputs, for each group of `Tuples`, a `Tuple` that represents/summarizes that group of `Tuples` (e.g., for each location code, the minimum temperature). An `AggregateOperator` is an operator that uses an `Aggregator` and supports the basic iterator capabilities (`hasNext()` and `next()` methods).

This part requires implementing only the following method:

- `createOutputSchema(AggregateFunction[] aggregateFunctions)`: constructs and then returns the output schema of the `AggregateOperator`. The output schema consists of the grouping attributes (e.g., `Location`) and additional attributes from `AggregateFunctions` (e.g., `Maximum(Temperature)`). This method needs to construct a new `RelationSchema` which requires an array storing the names of the attributes and another array storing the types of these attributes.

For the  $i$ -th grouping attribute, the name of that attribute can be obtained from `groupingAttributeNames[i]`. Given that attribute name, the type of the attribute can be obtained by using `inputSchema.attributeIndex(String)` and `inputSchema.attributeType(int)`. For the  $i$ -th `AggregateFunction`, the name of that `AggregateFunction` can be obtained from `aggregateFunctions[i].toString()`. The type of that `AggregateFunction` can be obtained from `aggregateFunctions[i].valueType()`.

When the above method is implemented correctly, your code will pass the unit tests in `AggregateOperatorTest` while showing the following output:

```
input schema: {ID=java.lang.Integer, Location=java.lang.Integer, Temperature=java.lang.Double}
grouping attributes: [Location]
aggregate functions: [Minimum(Temperature)]
output schema: {Location=java.lang.Integer, Minimum(Temperature)=java.lang.Double}
```

The above output shows a situation where (i) the input schema has attributes `ID`, `Location`, and `Temperature`, (ii) the grouping attribute is `Location`, (iii) the aggregate function is `Minimum(Temperature)`, and thus (iv) the output schema has attributes `Location` and `Minimum(Temperature)`.

#### Part 4. Aggregation: `Aggregator.java` (10 points)

As mentioned in Part 3, an `Aggregator` groups all `Tuples` (by user-specified attributes) from an input `Operator` and outputs, for each group of `Tuples`, a `Tuple` that represents/summarizes that group of `Tuples` (e.g., for each location code, the minimum temperature). In this part, you need to implement the following constructor and method in `Aggregator.java`:

- `Aggregator(Operator input, RelationSchema outputSchema, String[] groupingAttributeNames, Class<?>[] aggregateFunctionTypes, String[] aggregationAttributeNames)`: constructs an `Aggregator`. Given an input tuple  $t$ , the `Aggregator` needs to extract the values of the grouping attributes (e.g., `Location` value 0) and then finds the `AggregateFunctions` for that combination of grouping values (e.g., `Minimum` and `Maximum` that have been applied to the `Temperature` attribute from all `Tuples` whose `Location` value is 0). Then, the `Aggregator` needs to update all of these `AggregateFunctions` based on tuple  $t$  (e.g., update the minimum value if the `Temperature` value of  $t$  is smaller than the previous minimum). The above implementation approach requires space linear in the number of distinct groups. For the purposes of this assignment, you do not need to worry about the situations where there are too many groups to fit into the memory.
- `iterator()`: returns an iterator over the output `Tuples` of the `Aggregator`. The details of these output `Tuples` are explained above (refer to the output schema of the `AggregateOperator`).

When the above constructor and method are implemented correctly, your code will pass the unit tests in `AggregatorTest`.

### Appendix A. Importing a Java Project

1. Start Eclipse. If Eclipse runs for the first time, it asks the user to choose the workspace location. You may use the default location.

2. In the menu bar, choose “File” and then “Import”. Next, select “General” and “Existing Projects into Workspace”. Then, click the “Browse” button and select the “hdb\_query.zip” file contained in this assignment package.
3. Once the project is imported, you can choose `TupleArrayReaderTest.java`, `SelectionOperatorTest.java`, `textttAggregateOperatorTest.java`, or `AggregatorTest.java` and then run the program.